

Consistency of Safety Decisions Across Sampling Configurations
(Llama 3.1 8B Instruct, N=876 prompts)

